

CS531 Programming Assignment 2: Towers of Corvallis

Michael Lam, Xu Hu
EECS, Oregon State University

October 24, 2012

Abstract

In this assignment we design, implement and discuss two different informed search algorithms and heuristics to solve the Towers of Corvallis, which is a variation of Towers of Hanoi.

1 Introduction

The Towers of Corvallis puzzle is a variation on the Towers of Hanoi puzzle. While similarly consisting of 3 pegs and n disks, the Corvallis variation allows any disk to go on top of any other disk. The goal is to find the smallest number of moves in getting from an initial state to the goal state, which is defined as the order 9876543210 on peg A for 10 disks and similarly for fewer disks.

We implement two informed search algorithms: A* and RBFS (recursive best-first search). As informed searches, we also implement two heuristics, one admissible and one non-admissible. For each algorithm and heuristic function, we evaluate the performance by testing across different number of disks and different initial states.

2 A* Search

We implement the A* search algorithm for finding the paths between initial states to a given goal state in the problem of Tower of Corvallis. It is summarized in algorithm 1.

Algorithm 1 A* Search

```
exploredSet =  $\emptyset$ 
frontier = [initialPath]
while number(explored) < NMAX do
  if frontier ==  $\emptyset$  then
    return FALSE
  end if
  path = frontier.pop()
  state = path[0]
  exploredSet.add(state)
  if state == goalState then
    return path
  end if
  for action in state.validActions() do
    for newState in action.results() do
      newPath = path + newState
      if ismember(frontier,newPath) == FALSE then
        frontier.push(newPath)
      end if
    end for
  end for
end while
```

For implementing the frontier, we use the priority queue, which is actually a heap data structure. We use a callback function $f(state) = g(state) + h(state)$ as the priority, where $g(state)$ is the length of the path and $h(state)$ is the heuristic for estimating the distance between the current state to the goal state. We will analyze several admissible and non-admissible heuristics in section 4.

3 Recursive Best-First Search

Recursive best-first search or RBFS works by storing an f-limit for each node. The algorithm uses the f-limit to decide which subtree of the problem tree to explore by considering the best and 2nd best (alternative) f-limits. In order to keep the search functional, RBFS also updates the f-values of each node during the search.

The advantage of RBFS over A* is that RBFS uses less memory. Whereas A* stores all of its explored nodes, RBFS will only keep relevant nodes in memory. However, the disadvantage of RBFS over A* is that RBFS could expand more nodes than A* due to redundancy. Since RBFS does not store all nodes explored, it can re-expand the same nodes and thereby increasing computation time.

The pseudocode is listed in algorithm 2.

Algorithm 2 RBFS Search

```

function RBFS(state, f-limit):
  if state is goal state then
    return solution
  end if
  successor = all children of state
  if successors is empty then
    return failure
  else
    for s in successors do
      s.f = max(s.g + s.h, state.f)
    end for
    while true do
      best = lowest f-value node in successors
      if best.f > f-limit then
        return failure, best.f
      end if
      alternative = second best f-value of any node in successors
      result, best.f = RBFS(best, min(f-limit, alternative))
      if result is not failure then
        return result
      end if
    end while
  end if

```

We will analyze several admissble and non-admissble heuristics in section 4.

4 Experiments

4.1 Heuristics

We observed that the length of a path is the steps moving from initial state to goal state. So a natural heuristic can be designed based on an optimistic estimation of how many steps need to turn a state to its final goal. The first admissible heuristic is as following:

$$h(s) = (\sum_i |s(1, i) - i|) + \text{numberDisks}(s(2)) + \text{numberDisks}(s(3))$$

Here, $s(i, j)$ is the j th disks in peg i . The intuition is that for a disk k in peg 1, it has to take $|k - i|$ steps to get to the place it should be, where i is its current location on the peg. For example, if the 9th disk is on the bottom, it has to take $9 - 0$ steps to get to the top. For disks on peg 2 and peg 3, they have to be moved from the current location to peg 1 in at least 1 step. This is in fact an admissible heuristic. A simple justification is given as follows: for any disk on peg 1, before it reaches its ideal position, there should be $|k - i|$ disks put underneath it, and for every such disk, it takes at least one step. Of course, a much trivial admissible heuristic can be used here, such as number of disks in the first peg, but for computational efficiency, we choose this admissible heuristic in our experiment.

For non-admissible heuristic, we simply enlarge the admissible heuristic by a factor of 2. This is also reasonable since before we move a disk to its ideal position, usually we have to move another disk out of that place, which makes the number of steps double. But obviously this is not true for all the cases. We can give a counterexample easily.

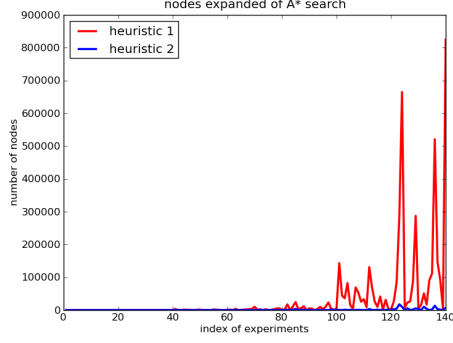
$$h(s) = 2 \times [(\sum_i |s(1, i) - i|) + \text{numberDisks}(s(2)) + \text{numberDisks}(s(3))]$$

4.2 Results

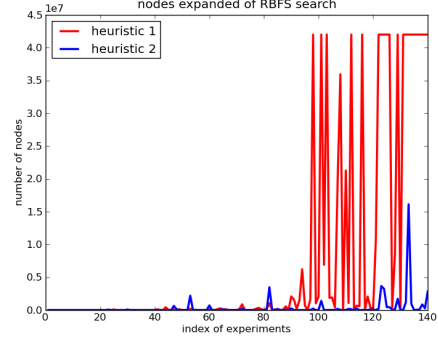
5 Discussion

We discuss our results and answer the questions in this section.

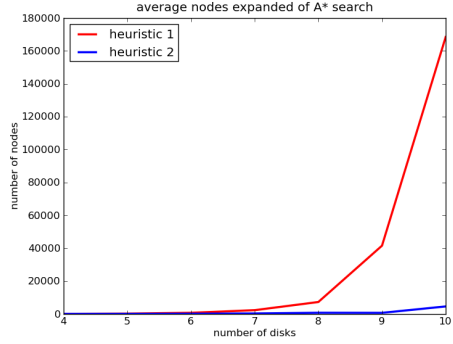
1. Show an example solution sequence for each algorithm for the largest size you tested




(a) A* nodes: all 140 experiments



(b) RBFS nodes: all 140 experiments



(c) A* average searched nodes



(d) RBFS average searched nodes


Figure 1: Plots for the number of nodes searched against the problem size for each algorithm and heuristic. For the first 2 figures, number of disks within 4-10, every 20 examples have the same problem size. We show average expanded nodes against number of disks for a clear demonstration. As the problem size becomes larger, the number of searched nodes grow up exponentially.

For size $n = 10$ and A* on problem "46910517328":


[46910517328,-,-], ...

For size $n = 10$ and RBFS on problem "7126049853":


[7126049853,-,-], [126049853,7,-], [26049853,17,-], [6049853,217,-], [049853,6217,-],
 [49853,06217,-], [9853,406217,-], [853,9406217,-], [53,89406217,-], [3,589406217,-],
 [-,3589406217,-], [3,589406217,-], [-,3589406217,-], [3,589406217,-], [-,3589406217,-],
 [3,589406217,-], [-,3589406217,-], [3,589406217,-], [53,89406217,-], [53,9406217,8],




(a) A*: average cpu time (s) on heuristics



(b) RBFS: average cpu time (s) on heuristics



(c) A*: average cpu time (s) on whole problem




(d) RBFS: average cpu time (s) on whole problem

Figure 2: Plots for the average cpu time against the problem size over 20 examples for each algorithm and heuristic.


[53,406217,98], [3,5406217,98], [3,406217,598], [3,06217,4598], [-,06217,34598],
 [0,6217,34598], [0,217,634598], [0,17,2634598], [10,7,2634598], [210,7,634598],
 [210,67,34598], [3210,67,4598], [43210,67,598], [543210,67,98], [6543210,7,98],
 [76543210,-,98], [76543210,9,8], [876543210,9,-], [9876543210,-,-]

2. Is there a clear preference ordering among the heuristics you tested considering the number of nodes searched and the total CPU time taken to solve the problems for the two algorithms?


The A* algorithm with nonadmissible heuristic performs the best among




(a) Nodes against disks for heuristic 1



(b) Nodes against disks for heuristic 2



(c) cpu time against disks for heuristic 1



(d) cpu time against disks for heuristic 2

Figure 3: Performance comparisons between A* and RBFS

Disk:	4	5	6	7	8	9	10
A*/admissible H1	8.8	11.6	13.95	16.5	19.0	n/a	n/a
A*/admissible H7	8.95	11.6	13.95	16.75	19.2	22.4	24.77
RBFS/admissible H9	8.8	11.6	13.95	16.35	17.11	n/a	n/a
A*/nonadmissible 1	8.85	11.85	14.6	17.5	20.2	24.1	27.6
RBFS/nonadmissible 1	8.85	11.85	15.25	19	22.21	26.81	32.25
A*/nonadmissible 2	9.25	12.95	16.5	20.05	23.8	27.4	34.05
RBFS/nonadmissible 2	11.25	17.85	24.35	32.75	40.35	50.7	64.55

Table 1: Average solution length per algorithm, heuristic and disk size. n/a means unable to compute within 10 minutes for all problems.

other heuristic type and algorithm combinations.

3. Can a small sacrifice in optimality give a large reduction in the number of nodes expanded? What about CPU time?

The non-admissible heuristic sacrifices optimality because it does not necessarily yield an optimal solution. However, the sacrifice in optimality can yield a large reduction in the number of nodes expanded and thereby reducing the overall CPU time. For instance, an admissible heuristic can severely underestimate the true cost and yield a large number of nodes expanded. However, even though a non-admissible heuristic overestimates the true cost, it could be the case that the overestimation is very little to the true cost. Therefore, there can be a large reduction in the number of nodes expanded and thus a reduction in CPU time, as compared to the admissible heuristic that severely underestimates the true cost.

4. How did you come up with your heuristic evaluation functions?

We created several heuristic evaluation functions based on some intuitions and for the sake of testing. One of our first heuristic evaluation functions is a simple count of the number of disks in the first peg. Then we gradually took advantage of the goal state's structure and compared it to the current state in our later heuristic designs. Hence, we came up with our heuristic evaluation functions by designing easier ones first, then gradually evolving them into more complex ones. We briefly describe all of our heuristics here, even if we only chose some for our experiments:

Heuristic 1 Number of disks in the first peg. RBFS.

Heuristic 2 Total number of disks MINUS the number of matches between goal state and first peg. RBFS.

Heuristic 3 Double the total number of disks MINUS the number of matches between goal state and first peg, AND MINUS the streak length of the first peg to the goal state starting at the bottom. RBFS.

Heuristic 4 Sum the distances of all current disk positions to their goal positions plus the lengths of the second and third pegs. A*.

Heuristic 5 Same as heuristic 4 for RBFS.

Heuristic 6 Same as heuristic 4, but enlarge heuristic value by factor of 2. A*.

Heuristic 7 Same as heuristic 4, but instead of summing the distances, take the max distance of every disk position to their goal position on the first peg. A*.

Heuristic 8 Same as heuristic 6 for RBFS.

Note: since our two algorithms utilize different data structures, some heuristics are equivalent but use different data structures. Our organization might not have been the best.

5. How do the two algorithms compare in the amount of search involved and the cpu-time?

According to our results, in general, A* does a better job by expanding less nodes than RBFS and using less cpu-time. Furthermore, the nonadmissible heuristic does a better job than the admissible heuristic. We can see these on our graphs, where RBFS has higher values of number of expanded nodes and cpu-time as the number of disks increases, and similarly for the admissible heuristic.

6. Do you think that either of these algorithms scale to even larger problems? What is the largest problem you could solve with the best algorithm+heuristic combination? Report the wall-clock time, CPU-time, and the number of nodes searched.

The algorithms could scale to even larger problems if the heuristic were designed very well. It is amazing to see an algorithm with one heuristic run very slowly with many node expansions while the same algorithm with a better heuristic finishing the job much faster with significantly fewer node expansions. Therefore, a carefully designed heuristic will allow both algorithms to scale to even larger problems by reducing the amount of expanded nodes and thereby reducing CPU time.

A smaller factor in scaling to larger problems is the algorithm design. A* may run out of memory eventually compared to RBFS since A* stores every node it expands, but this would take awhile since our computers contain a lot of memory. RBFS could scale further since it does not store every node in memory, but it will take longer to finish the job due to that trade-off. However, changing the heuristic would have the most impact because the

overall number of expanded nodes would be reduced with a better heuristic, regardless of the algorithm chosen.

The largest problem we could solve is of disk size 10, using A* and the non-admissible. It took less than ten minutes. For the initial state of "", it was able to solve the problem in 1.41 seconds and 6536 expanded nodes, with a solution length of 34 steps. As a non-admissible heuristic, it may not necessarily be the optimal solution, but the problem was nonetheless solved.

We went further and tested disks sizes above 10. We tested up to size 15. Given an initial state of

0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14

, the A* algorithm using the non-admissible heuristic was able to solve the problem in 2687 seconds and 987156 expanded nodes, with a solution length of 62 steps. That is about 44 minutes. We believe we can go higher, but it will take hours to complete. This demonstrates that our algorithms and heuristics can scale to even larger problems.

7. Is there any tradeoff between how good a heuristic is in cutting down the number of nodes and how long it took to compute? Can you quantify it?

A heuristic function $H1$ can take longer to compute than heuristic function $H2$, but if $H1$ significantly reduces the number of nodes, then the overall time to solve the problem could be less using $H1$ than $H2$. In every experiment, we measured every time it takes to compute a heuristic function and averaged them. To take one concrete example, the heuristic of counting the number of disks in the first peg takes shorter computation time than the heuristic of calculating matches between the first peg and the goal state. As the disk size increases, the computation time noticeably but slightly increases. However, the matching heuristic performs better than the counting heuristic because it reduces the number of expanded nodes significantly. Therefore, while a heuristic evaluation function may take longer to compute, it may reduce the overall number of expanded nodes to make a significant reduction in overall computation time of the problem.

8. Is there anything else you found that is of interest?

We observe that our admissible heuristic function yields different solution lengths, ever so slightly, for the different algorithms. We do not expect this,

since a heuristic function should yield an optimal solution and thus have the same lengths. That either suggests that our admissible heuristic function is not indeed admissible, or that there is a problem in our heuristic or algorithm, or computation results.