

CS531 Programming Assignment 4: Wumpus Agent

Michael Lam, Xu Hu
EECS, Oregon State University

November 29, 2012

Abstract

In this assignment we design, implement and evaluate an algorithm that uses first-order logic and A* search for an agent in order to solve Wumpus puzzles.

1 Introduction

The Wumpus world is a 4x4 grid containing pits, one Wumpus and one gold at various locations. The objective of the agent is to retrieve the gold without dying from the Wumpus or falling into a pit. Furthermore the agent can perceive its environment and infer the location of pits and the Wumpus due to the rules of constructing a Wumpus world. Therefore it makes sense to implement an algorithm involving logic and search to make intelligent decisions.

We used the existing Wumpus environment simulator for Python provided by Walker Orr. We designed an agent that uses first-order logic with a tell-ask interface to assert/query what it knows about the Wumpus world and A* search to plan routes around the Wumpus world. The algorithm is essentially the same as the one provided in the Russell-Norvig textbook (pg. 270, fig. 7.20). For answering logic queries, the program uses the Prover-9 program.

To evaluate our algorithm, we collected statistics on the agent such as if the agent were successful. We ran experiments on 40 randomly generated maps (though some are unsolvable with gold in the same square as a pit).

We also evaluated our agent using an RBFS algorithm instead of A* and compared results.

2 Approaches

In this section we describe our implementation of the agent and discuss our design decisions.

2.1 Simulator

We used the simulator for Python provided by Walker Orr. While the simulator represented the Wumpus world fairly well as described in the textbook, it was still incomplete. First of all, the simulator lacked the notion of facing directions as well as the percepts bump and scream. In addition, the simulator would declare reaching the gold a success, eliminating the need for the agent to travel back and climb at the initial square. This also differs from the textbook's formulation. Therefore, we modified the simulator to add these remaining percepts, fixed the simulator to keep track of facing direction and allowed the agent to return to the initial square to climb. The modified simulator should represent the Wumpus world as described in the textbook.

2.2 Knowledge Base

We designed our Knowledge Base as follows:

- Atemporal

B(x,y) Breezy at location (x,y)

P(x,y) Pit at location (x,y)

S(x,y) Smelly at location (x,y)

W(x,y) Wumpus at location (x,y)

- Temporal

WumpusAlive(t) Wumpus is alive at time t

HaveGold(t) Agent has gold at time t

HaveArrow(t) Agent has arrow at time t

- Both

Loc(x,y,t) Agent was at location (x,y) at time t

OK(x,y,t) Square (x,y) at time t was safe to visit

- Perception

Breeze(t) Breeze perceived at time t

Stench(t) Stench perceived at time t

Glitter(t) Glitter perceived at time t

Bump(t) Bump perceived at time t

Scream(t) Scream perceived at time t

- Action

Forward(t) Forward action at time t

TurnLeft(t) TurnLeft action at time t

TurnRight(t) TurnRight action at time t

Shoot(t) Shoot action at time t

Grab(t) Grab action at time t

Climb(t) Climb action at time t

We wrote out some rules that Prover-9 could use to prove a query. Note that in our implementation, we explicitly fill for each x , y and t instead of leaving it with variables. We know the size of the world is 4x4 so we are able to enumerate all x and y . For each time step, we enumerate all x and y for time t . The reason for enumerating out all possibilities instead of using variables is to keep the logic relatively simple and avoid having to implement arithmetic in Prover-9.

- $B(x, y) \Leftrightarrow P(x1, y1) \vee P(x2, y2) \vee P(x3, y3) \vee P(x4, y4)$
- $S(x, y) \Leftrightarrow W(x1, y1) \vee W(x2, y2) \vee W(x3, y3) \vee W(x4, y4)$
- $Loc(x, y, t) \Rightarrow (Breeze(t) \Leftrightarrow B(x, y))$
- $Loc(x, y, t) \Rightarrow \neg P(x, y)$

- $Loc(x, y, t) \Rightarrow (Stench(t) \Leftrightarrow S(x, y))$
- $Loc(x, y, t) \Rightarrow (\neg W(x, y)) \vee (W(x, y) \wedge \neg WumpusAlive(t))$
- $OK(x, y, t) \Leftrightarrow \neg P(x, y) \wedge \neg(W(x, y) \wedge WumpusAlive(t))$

Other than these fundamental rules, the rest are essentially assertions of percepts, actions and derived atemporal facts of the Wumpus world at every time step. This information should be sufficient for the agent to reason about the world and make good decisions.

2.3 Hybrid Agent

We implemented the algorithm in the textbook (pg. 270, fig. 7.20). It is mostly the same in terms of the decision-making process. There are a few technical differences regarding the Knowledge Base representation.

We asserted all the facts at the beginning of the algorithm. The Make-World-Logic-Sentences function constructs sentences that assert conditional sentences of the Wumpus world at the current time. This includes sentences relating "OK" to "Pit" and "Wumpus" as well as explicitly writing out the sentences relating adjacent squares such as relating "Breeze" to "Pit." (Rules listed out in the previous section.) It is important to note that some of these sentences are temporal so we create explicit rules at every new time step rather than a general rule for all time steps. Again, this is to avoid adding arithmetic logic to Prover-9. Another note is that the adjacent squares are written out explicitly for the same reason.

One final technical note is that every call to Ask also caches the query if the query is proven true. This is to improve performance.

Algorithm 1 Hybrid-Wumpus-Agent

HYBRID-WUMPUS-AGENT(percept)

Inputs : "percept" list

Persistent : KnowledgeBase "KB", time "t", actionsequence "plan"

Returns : singlenext "action"

Tell(KB, Make-World-Logic-Sentences(t))

Tell(KB, Make-Percept-Sentence(percept, t))

Tell(KB, Make-Location-Safe-Sentence(current))

Tell(KB, Make-Have-Arrow-Sentence())

Tell(KB, Make-Have-Gold-Sentence())

safe := {(x,y) : Ask(KB, OK(x,y,t)) = TRUE}

if Ask(KB, Glitter(t)) = TRUE **then**

 plan := [Grab] + Plan-Route(current, {(0,0)}, safe) + [Climb]

end if

if plan is empty **then**

 unvisited := {(x,y) : ASK(KB, Loc(x,y,t')) = FALSE for all t' <= t}

 plan := Plan-Route(current, unvisited \cap safe, safe)

end if

if plan is empty AND Ask(KB, HaveArrow(t)) = TRUE **then**

 possible_wumpus := {(x,y) : Ask(KB, -W(x,y)) = FALSE}

 plan := Plan-Shot(current, possible_wumpus, safe)

end if

if plan is empty **then**

 not_unsafe := {(x,y) : Ask(KB, -OK(x,y,t)) = FALSE}

 plan := Plan-Route(current, unvisited \cap not_unsafe, safe)

end if

if plan is empty **then**

 plan := Plan-Route(current, {(0,0)}, safe) + [Climb]

end if

action := Pop(plan)

Tell(KB, Make-Action-Sentence(action, t))

t := t+1

return action

2.4 Route Planning

In the Plan-Route function we formulate a search problem given the agent’s current location, goal locations and allowable squares. We decided to use A* as the search algorithm but we also implement RBFS to compare and evaluate the performance of the two algorithms. We implement the A* search and RBFS searches such that they are essentially the same as in the previous assignment but with a different heuristic. The state space is defined by position and direction, that is a tuple $(x, y, direction)$. Since the agent may have multiple goals, the heuristic is the minimum city block distance from the current square to one of the goals. When the path is found, we simply backtrack the squares on the path and return a sequence of actions, which will be combined with other actions in hybrid agent aforementioned.

3 Experiments

We tested our agent on 40 randomly generated maps and used A* search and RBFS search respectively. For running time considerations, we only tested with 4×4 maps. All other conditions are the same including heuristics and the logic engine. We record 5 targets: returning successfully ("success"), grabbing gold ("gold"), killing wumpus ("kill"), number of steps ("steps") and points of reward ("reward"). For rewards, we set every action to -1 , as well as 500 for grabbing the gold, 100 for killing the wumpus and 200 for returning back and climbing out successfully.

The details for each experiment are displayed in table 1. Note that there are many unsolvable maps. For A* search, the agent finished 17 games, got 17 golds and killed the wumpus 27 times. The average steps and rewards for A* search are 119.75 and 210.25. The RBFS search couldn’t finish some of searches because we set a relative small number of allowed expanded nodes. It completed 10 games, got 17 golds, killed wumpus 27 times. The average steps for RBFS search is 96.55, while the average reward is 198.45. From these we can see that the performances of two search algorithms are similar. A* search is slightly better.

4 Discussion

From the results, it appears that in some circumstances the agent is able to climb back out without grabbing any gold. This is good because the net reward in the end could be positive rather than negative. However, we still count success as retrieving gold and climbing back out so coming out without gold is still not ideal.

On the other hand, the agent still fails at some maps even if they are solvable. That is probably because the agent had to take a risk step to explore as indicated in one if-branch of the algorithm. Therefore we cannot design a perfect agent unless the environment is fully observable and the agent can know everything in the beginning.

We initially implemented first-order logic sentences such that a sentence contains many quantified variables. While this one sentence captures everything, it is also ridiculously slow during execution time because Prover-9 had to convert the sentence into a form suitable for proving. The other problem with using quantified temporal variables is the need for arithmetic. Since Prover-9 does not recognize arithmetic out of the box, we would have to implement our own predicates for arithmetic. To circumvent these issues, we enumerated sentences over all possible quantified variable values. It turns out that switching to this methodology dramatically increased performance.

Another performance issue with Prover-9 is the need to convert percepts to nonfluent facts as soon as possible. Initially we told the knowledge base sentences that would derive nonfluent facts from fluent facts. In fact almost all the "calculation rules" were asserted into the knowledge base. Again, this was extremely slow. Therefore we decided to "compute" these facts before asserting facts into the knowledge base. By only asserting facts into the knowledge base for these, our performance improved dramatically once again.

In future experiments when we would have more time, we would try out for world sizes other than 4×4 . It would be interesting to see at what point the algorithm would slow down so much as to make the problem intractable. In fact, our algorithm gradually slows down as the knowledge base gets larger. Another interesting experiment would be to isolate out map classes, such as maps that are unsolvable because there is no path to the gold. Would the agent make exploration risk moves and die, or climb out without gold?

A*	Index	1	2	3	4	5	6	7	8	9	10
	Solvable	1	1	0	1	1	1	0	0	1	0
	Success	1	0	0	1	0	0	0	0	0	0
	Gold	1	0	0	1	0	0	0	0	0	0
	Kill	0	0	1	0	0	0	0	1	0	1
	Steps	241	12	14	251	329	315	230	60	175	192
	Reward	559	88	-14	449	-229	-215	-130	-60	-75	-92
RBFS	Index	1	2	3	4	5	6	7	8	9	10
	Solvable	1	1	0	1	1	1	0	1	1	1
	Success	0	0	0	0	1	1	0	1	0	1
	Gold	0	0	0	0	1	1	0	1	0	1
	Kill	0	0	0	0	1	1	1	1	0	1
	Steps	205	60	67	7	99	98	7	39	211	146
	Reward	-105	-60	-67	-7	601	602	-7	661	-111	554
A*	Index	11	12	13	14	15	16	17	18	19	20
	Solvable	0	0	1	1	1	0	1	1	1	1
	Success	0	0	1	1	1	0	0	1	1	1
	Gold	0	0	1	1	1	0	0	1	1	1
	Kill	0	0	1	1	1	1	1	1	1	1
	Steps	65	7	99	39	99	90	7	162	77	235
	Reward	35	-7	601	661	601	-90	-7	538	623	465
RBFS	Index	11	12	13	14	15	16	17	18	19	20
	Solvable	0	0	1	1	1	0	1	1	1	1
	Success	0	0	1	0	0	0	1	0	0	1
	Gold	0	1	1	0	1	0	1	0	0	1
	Kill	0	1	1	0	1	0	1	0	1	1
	Steps	19	168	130	14	174	7	122	60	161	99
	Reward	81	332	570	-14	326	-7	578	-60	-61	601
A*	Index	21	22	23	24	25	26	27	28	29	30
	Solvable	1	1	1	1	1	1	1	0	1	1
	Success	1	0	0	1	1	0	1	0	0	0
	Gold	1	0	0	1	1	0	1	0	0	0
	Kill	1	0	0	1	1	0	0	0	0	0
	Steps	99	7	7	175	122	19	271	67	7	206
	Reward	601	-7	-7	525	578	81	529	-67	-7	-106
RBFS	Index	21	22	23	24	25	26	27	28	29	30
	Solvable	1	1	1	1	1	1	0	0	1	1
	Success	0	0	0	0	1	0	0	0	0	1
	Gold	0	1	0	0	1	1	0	0	0	1
	Kill	0	1	0	1	0	1	0	0	0	1
	Steps	12	159	138	90	39	152	72	7	65	77
	Reward	88	341	-38	-90	661	348	28	-7	35	623
A*	Index	31	32	33	34	35	36	37	38	39	40
	Solvable	1	1	1	1	1	1	1	1	1	0
	Success	1	1	0	0	0	0	1	1	1	0
	Gold	1	1	0	0	0	0	1	1	1	0
	Kill	1	1	0	0	0	0	1	1	1	0
	Steps	235	98	72	60	7	60	39	174	294	72
	Reward	465	602	28	-60	-7	-60	661	526	406	28
RBFS	Index	31	32	33	34	35	36	37	38	39	40
	Solvable	1	1	1	1	1	1	1	1	0	0
	Success	0	0	1	0	0	0	0	0	0	0
	Gold	0	0	1	0	1	1	0	0	1	0
	Kill	0	0	1	0	1	0	1	0	0	0
	Steps	260	214	99	7	158	193	60	72	88	7
	Reward	-160	-114	601	-7	342	407	-60	28	512	-7

Table 1: Statistics of all 40 experiments, one with A* and one with RBFS. Solvable indicates that the map is solvable, which is only defined as if the gold is not in the same square as a pit.